

# S3 Performance Analysis Report

## Group Members:

- Clement Owusu Bempah
- Marie-Pearl Otoo
- Michael Aboagye
- Samuel Kofi Asare Dwumah
- Vera Sekyere

## Contents

S3 Performance Analysis Report .....	1
Group Members: .....	1
Table of Contents .....	<b>Error! Bookmark not defined.</b>
Introduction .....	1
1. Transfer Acceleration Review .....	1
2. Activity Metrics Analysis .....	4
3. Status Code Metrics Review .....	4
4. Recommendations for S3 Performance Improvement .....	6

## Introduction

This report summarizes key findings from an analysis of Amazon S3 performance, focusing on metrics, tools, and features used, and provides actionable recommendations for improvement. The analysis covers S3 Transfer Acceleration, activity metrics (GetRequests, PutRequests, BytesDownloaded), and status code metrics (200OKStatusCount, 403ForbiddenErrorCount, 404NotFoundErrorCode).

## 1. Transfer Acceleration Review

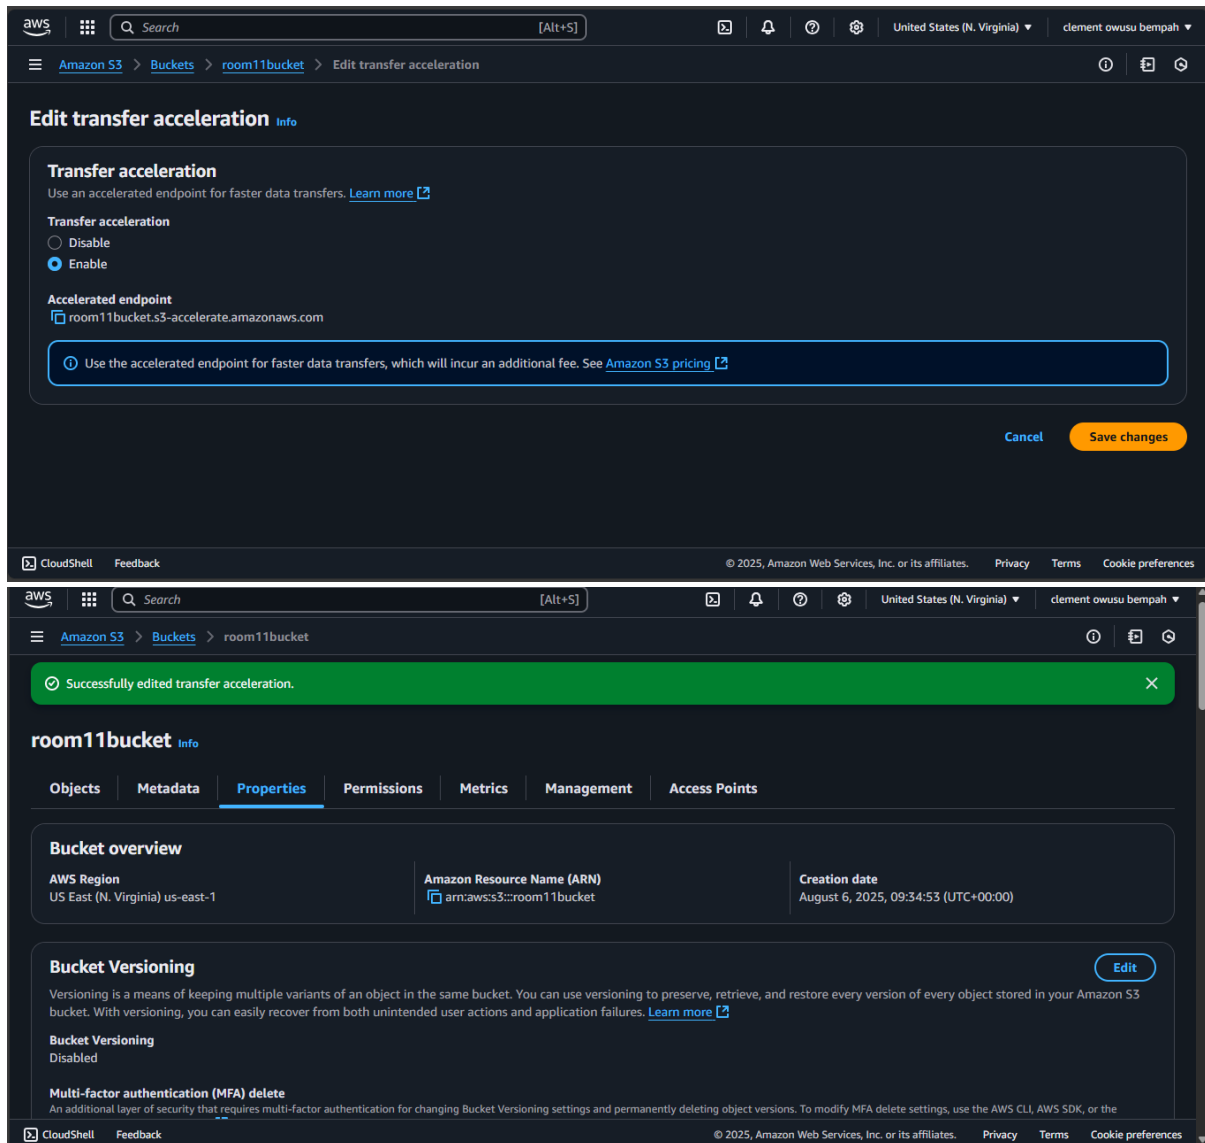
**S3 Transfer Acceleration** leverages Amazon CloudFront's globally distributed edge locations to optimize data transfers to and from S3 buckets. When enabled, data is routed through the nearest edge location, which then uses optimized network paths to transfer the data to the S3 bucket. This significantly improves upload and download speeds, especially over long geographical distances, by reducing latency and leveraging Amazon's high-speed backbone network.

Enabling Transfer Acceleration:

As demonstrated in the provided screenshots, Transfer Acceleration can be enabled directly from the S3 bucket properties in the AWS Management Console. Navigate to the bucket, select the "Properties" tab, and then locate the "Transfer acceleration"

section to enable it.

Alternatively, it can be configured programmatically via the AWS CLI or SDKs using commands like `aws configure set default.s3.use_accelerate_endpoint true` for uploads and downloads.



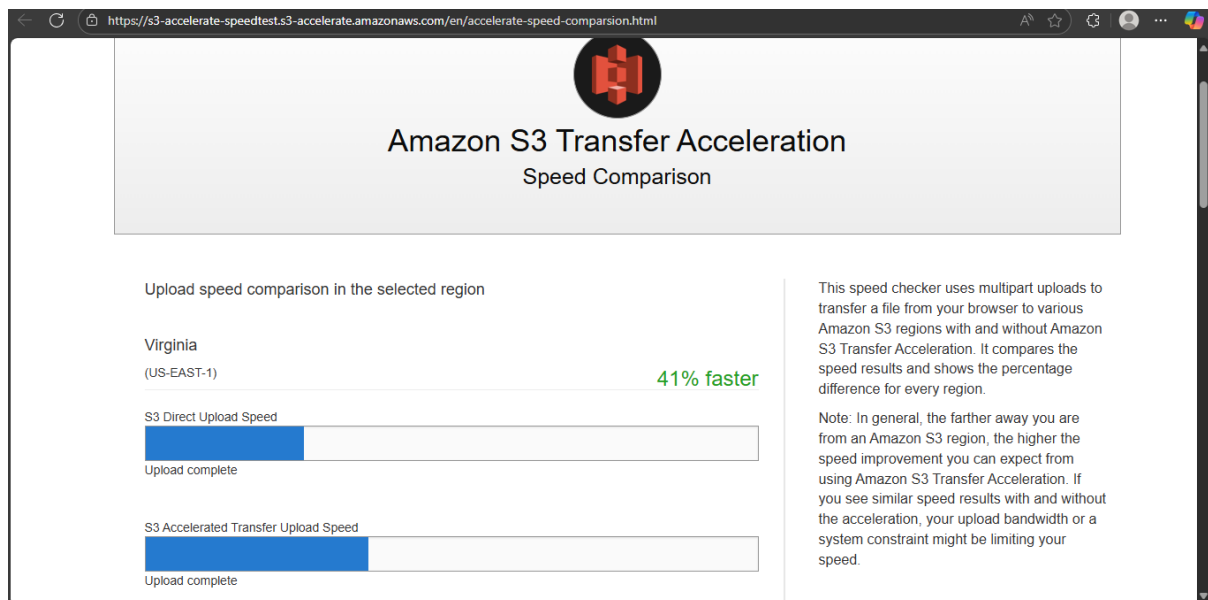
## Pros:

- **Faster Transfers:** Substantially improves data transfer speeds, particularly for users geographically distant from the S3 bucket's region.
- **Improved User Experience:** Reduces wait times for large file uploads and downloads, leading to a better experience for end-users.
- **Global Reach:** Benefits applications with a globally dispersed user base.

## Cons:

- **Additional Cost:** Transfer Acceleration incurs an additional fee, which is based on the amount of data transferred through the accelerated endpoint.
- **Not Always Beneficial:** For users located close to the S3 bucket's region, the performance improvement might be minimal and not justify the extra cost. The speed comparison tool showed a 41% faster upload in Virginia (US-EAST-1), indicating its effectiveness in certain scenarios.

```
cleme@DESKTOP-BHD3JTJ MINGW64 ~/Downloads
$ aws s3 cp C:\\Users\\clleme\\Downloads\\html5up-ethereal.zip s3://room11bucket/
upload: .\\html5up-ethereal.zip to s3://room11bucket/html5up-ethereal.zip
```

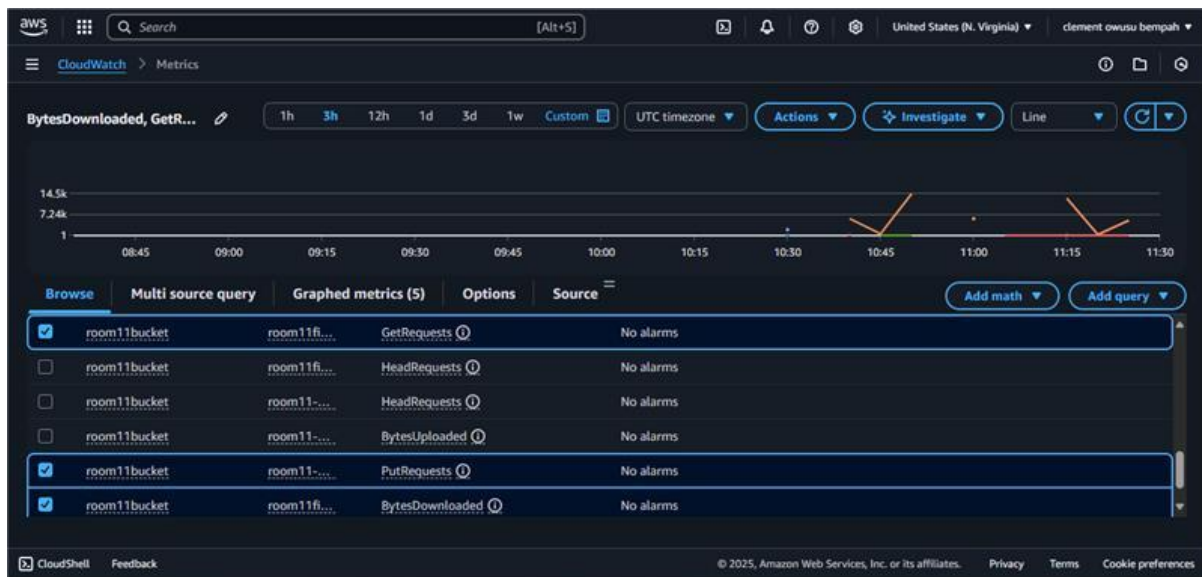


```
cleme@DESKTOP-BHD3JTJ MINGW64 ~/Downloads
$ aws s3 ls s3://room11bucket/
2025-08-06 11:26:53      1762389 html5up-ethereal.zip
2025-08-06 10:42:12      232443  image (10).png
2025-08-06 10:43:14      94875  image (2).png
2025-08-06 10:32:52      46474  image (28).png
2025-08-06 10:32:51      66567  image (29).png
2025-08-06 10:43:10      157581 image (3).png
2025-08-06 10:32:49      63751  image (30).png
2025-08-06 10:32:48      19217  image (4).jpg
2025-08-06 10:42:56      161151 image (4).png
2025-08-06 10:42:52      157581 image (5).png
2025-08-06 10:42:46      11298  image (6).png
2025-08-06 10:42:45      165394 image (7).png
2025-08-06 10:42:27      144251 image (8).png
2025-08-06 10:42:18      237436 image (9).png
2025-08-06 11:16:37       31799 scaling.png
2025-08-06 11:15:02       31799 test.png
```

```
cleme@DESKTOP-BHD3JTJ MINGW64 ~/Downloads
$ aws s3 cp s3://room11bucket/test.png .
download: s3://room11bucket/test.png to .\test.png
```

## 2. Activity Metrics Analysis

Analyzing activity metrics like GetRequests, PutRequests, and BytesDownloaded provides insights into how an S3 bucket is being utilized and helps identify potential performance bottlenecks. The CloudWatch metrics dashboard shows the ability to monitor these metrics.

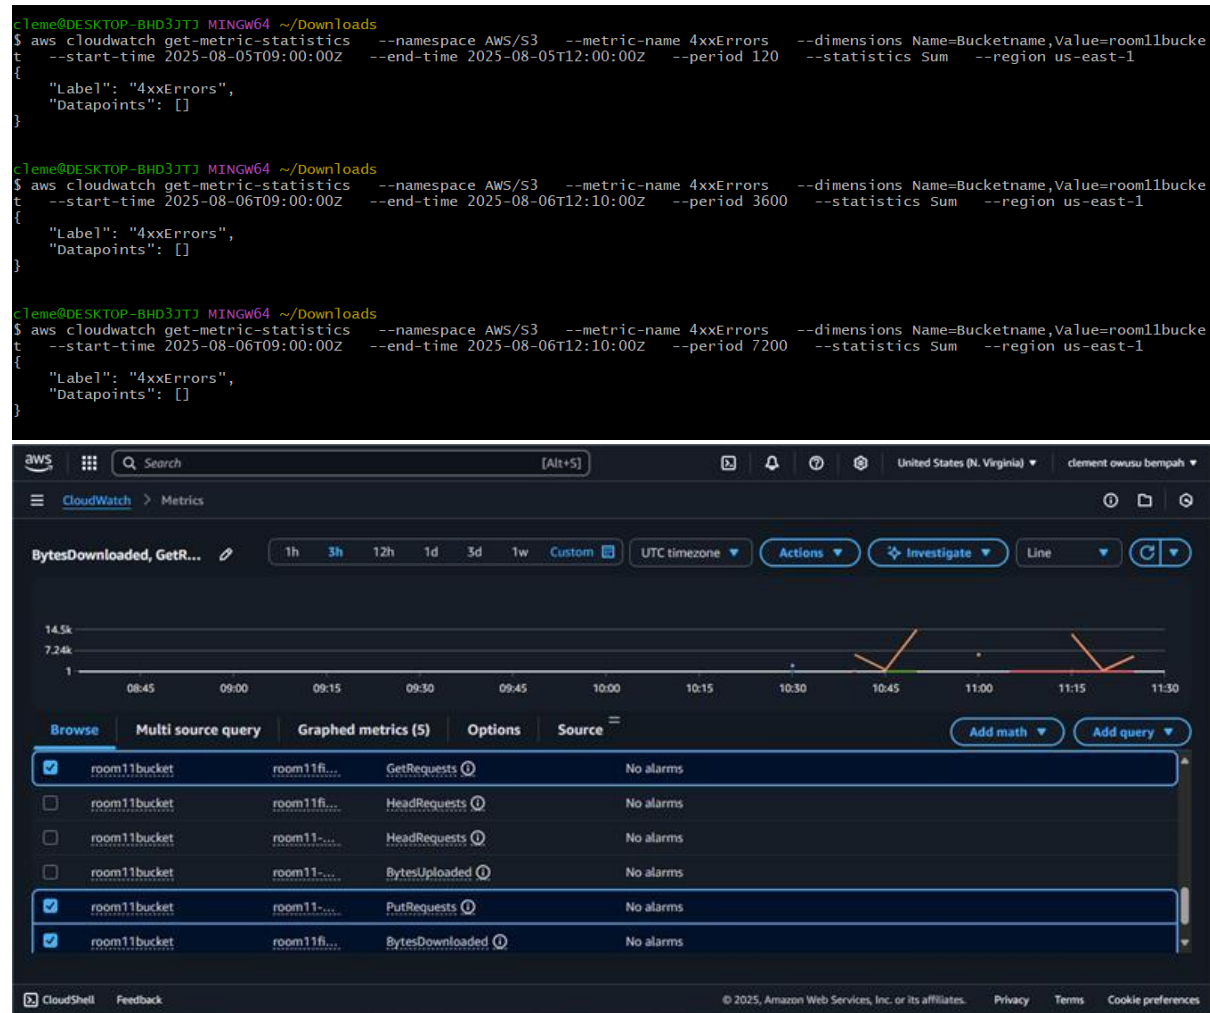


- **GetRequests:** Represents the number of times objects are retrieved from the bucket. A high volume of GetRequests could indicate frequent data consumption by applications or users.
- **PutRequests:** Indicates the number of times objects are uploaded or modified in the bucket. A high number here suggests frequent data ingestion or updates.
- **BytesDownloaded:** Shows the total bytes transferred out of the bucket. This metric is crucial for understanding data egress patterns and potential costs associated with data transfer out.

Based on the provided CloudWatch screenshot, GetRequests, PutRequests, and BytesDownloaded are being monitored. Without specific data points or trends over time, it's difficult to pinpoint exact bottlenecks. However, consistently high spikes in GetRequests or PutRequests might suggest a need for scaling application components that interact with S3 or optimizing data access patterns. For example, if GetRequests are consistently high from a single geographical region, enabling Transfer Acceleration for that region would be beneficial.

## 3. Status Code Metrics Review

Monitoring status code metrics such as 200OKStatusCount, 403ForbiddenErrorCount, and 404NotFoundErrorCode is essential for understanding the success rate of requests and troubleshooting failed operations. The CloudWatch dashboard shows 4xxErrors being monitored.



- **200OKStatusCount:** Indicates successful requests. A consistently high count is desirable.
- **403ForbiddenErrorCount:** Signifies requests that were denied due to insufficient permissions.
- **404NotFoundErrorCode:** Represents requests for objects that do not exist in the bucket.

### Causes of Failed Requests and Troubleshooting:

- **403ForbiddenErrorCount:**
  - **Causes:** Incorrect bucket policies, IAM user/role permissions, or ACLs. Public access blocks might also be enabled, preventing unauthorized access.
  - **Troubleshooting:**
    - Review bucket policies and IAM permissions attached to the users or roles

attempting to access the bucket.

- Check object ACLs if specific objects have different permissions.
  - Verify if S3 Public Access Block settings are too restrictive for the intended access.
  - Use AWS CloudTrail to log API calls and identify the exact permission issue.
- **404NotFoundErrorCode:**
    - **Causes:** Incorrect object keys (filenames), objects being deleted, or requests being made to the wrong bucket.
    - **Troubleshooting:**
      - Verify the object key (path and filename) in the application or request.
      - Check if the object was accidentally deleted. S3 Versioning (if enabled) could help recover deleted objects.
      - Ensure the request is targeting the correct S3 bucket and region.

The provided CloudWatch screenshot for 4xxErrors shows some activity around 11:15, indicating a few client-side errors occurred. Further investigation into the specific 403 or 404 counts during that period would be necessary to diagnose the exact cause.

## 4. Recommendations for S3 Performance Improvement

Based on the analysis, here are three actionable steps to improve S3 performance:

1. **Optimize for Geographic Distribution with Transfer Acceleration:**
  - **Actionable Step:** Enable S3 Transfer Acceleration for buckets that serve users or applications geographically distant from the bucket's region. Continuously monitor the cost implications and actual performance gains to ensure it provides a tangible benefit.
  - **Best Practice:** Use the S3 Transfer Acceleration Speed Comparison tool to evaluate potential gains before committing to its use for all traffic. For applications with a global user base, consider using Amazon CloudFront as a content delivery network (CDN) in conjunction with S3, as it offers more advanced caching and edge delivery capabilities, often at a lower cost for highly accessed content.
2. **Implement Object Lifecycle Management and Tiering:**
  - **Actionable Step:** Configure S3 Lifecycle rules to automatically transition objects to more cost-effective storage classes (S3 Standard-IA, S3 Glacier) based on access patterns and age. For frequently accessed data, ensure it remains in S3 Standard.
  - **Best Practice:** Analyze GetRequests and BytesDownloaded metrics over time to identify data that is accessed infrequently. For example, if data hasn't been accessed in 30 days, transition it to S3 Standard-IA. This optimizes storage

costs without significantly impacting performance for active data.

### 3. **Leverage S3 Event Notifications for Real-time Processing:**

- **Actionable Step:** For applications that need to react to S3 object changes (new file uploads), configure S3 Event Notifications to trigger AWS Lambda functions. This allows for immediate processing of new data without polling the bucket, reducing latency and improving overall system responsiveness.
- **Best Practice:** Use S3 Event Notifications for tasks like image resizing, data validation, or initiating data pipelines upon object creation or modification. This eliminates the need for applications to constantly check the bucket for new objects, reducing GetRequests and improving efficiency.

These recommendations combine configuration changes within S3 and AWS best practices to enhance performance, optimize costs, and improve the reliability of S3-based applications.